

# **Kubernetes for beginners**

## **Kubernetes Overview (S2)**

### **Containers**

- eg. Docker
- containerize applications
- isolated environments with own processes and own network etc.
- like VM's but share the same kernel
  - different linux distributions are possible to run
  - Containers don't include the OS
- Containerized applications on dockerhub
- Docker run commands to start containers
- Containers are instances of images

### **Container Orchestration**

- Automatically deploying and managing containers
- docker swarm, kubernetes, mesos
- multiple instances easy to deploy as a service level

### **Kubernetes Architecture**

- Nodes: physical or virtual machine kubernetes is running on
- Cluster: set of nodes running on a cluster
- Master: node with kubernetes configured as master
  - Watches over other (worker) nodes and balances loads handles crashes, ...
- Kubernetes contains:
  - API Server (master)
  - etcd (master)
  - kubelet (worker)
  - Container Runtime (worker, eg. Docker)
  - controller (master)
  - scheduler (master)
- kubectl
  - command line tool for kubernetes

## **Setup Kubernetes (S3)**

- local (minikube, kubeadm), cloud(aws, gcp, azure), play-with-k8s.com

### Minikube

- preconfigured single node cluster

### Kubeadm

- <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>
- setup multi-node cluster on separate machines
- install container runtime on all nodes
- install kubeadm tool on all nodes
- initialize master
  - <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/create-cluster-kubeadm/#pod-network>
  - For problems with DNS:
    - `sudo nano /etc/NetworkManager/NetworkManager.conf`
    - comment out `#dns=dnsmasq`
    - `sudo service network-manager restart`
    - `kubectl -n kube-system delete pod -l k8s-app=kube-dns`
- join worker nodes with command given by initializing master node

### Google Cloud Platform

- go to kubernetes tab
  - click create cluster
  - connect cluster when it's initialized

<https://labs.play-with-k8s.com/> → simple test environment for kubernetes

## **Kubernetes Concepts (S4)**

### **POD's**

- <https://kubernetes.io/docs/concepts/workloads/pods/pod-overview/>
- single instance of application with a docker container running inside
  - 1 container per POD usually
  - Or helper container within the same POD
- creating new pods to scale up/ or adding new nodes
- deploy pod with *kubectl run <application> --image <image-name>*
  - downloads image from dockerhub or private repo
- *kubectl get pods* to see all pods
- *kubectl describe pod <podname>* to see information of pod

### **YAML Introduction (S5)**

- <https://learn.getgrav.org/16/advanced/yaml>

## Kubernetes Concepts PODs, ReplicaSets, Deployments (S6)

- yamls as inputs for creation of objects
- *kubectl create -f ...yaml*

### Replication Controllers and ReplicaSets

- controllers: monitor objects and respond
- Replication Controller →
  - ensures a defined number of instances run at all times
  - load balancing between pods across nodes
  - defined with yaml
  - *kubectl get replicationcontroller*
- Replica Sets
  - New way of replication controller
    - Api version is different
    - Requires selector with matchLabels to specify the pods
      - Monitors pods with the specified labels
    - *kubectl get replicaset*
    - scale replica set:
      - *kubectl replace -f <replica-definition file>*
      - *kubectl scale --replicas=<number> <replica-definition-file>*

### Deployments

- allows for rolling updates, rollbacks, pause update, ...
  - deployment object over replicaset
  - deployment-definition.yml similar to replicaset
  - *kubectl get deployments*
  - *kubectl get all*
  - in productive mostly just creating deployment definition file
- 
- new deployment triggers new rollout which creates new revision
    - creating also a new replica set
  - *kubectl rollout status <deployment-name>*
  - *kubectl rollout history <deployment-name>*
  - deployment strategies:
    - 1. Recreate: destroy older versions and then deploy newer
    - 2. Rolling Update(default): Take down and bring up one by one
  - Change deployment-definition yaml *and kubectl apply -f <deployment-name>*
  - Rollback to go to previous revision
    - *kubectl rollout undo <deployment-name>*

### **Networking in Kubernetes (S7)**

- every POD has its own IP address
  - down over local network in Node
- every Node has its own IP address
  - when cluster is setup Networking is not setup automatically
  - network has to be setup manual
    - use calico, flannel, weave-net, nsx ... to do it
    - will create a network with routing techs

## Services (S8)

- enable communication between components of application
- enable connectivity between applications/users

### NodePort Service

- access from outside node done with service
- forward requests on Node Ports to Port on POD Port
- acts like virtual server inside the node (has own IP address)
- create with definition file
  - set ports in spec definition (targetPort, port, nodePort)
  - selectors to connect with specific POD
- *kubectl create <service-definition>*
- *kubectl get services*
- automatically selects POD for LoadBalancing
  - service spans across nodes

### ClusterIp Service

- Creates virtual IP inside Cluster
- create with definition file
  - set ports in spec definition (targetPort, port)
  - selectors to connect with specific POD
- *kubectl create <service-definition>*
- *kubectl get services*

### LoadBalancer Service

- Distribute load across servers
- Allows access from outside